

Checkmatr Written Report

Andrew Mascillaro, Nabih Estefan

Overview

Chess has simple rules, but the game space is extremely large. With the strength of our current computers, we aren't even close to solving the game of chess by brute force. As a fallback, engines use a heuristic to try to find the best moves for a given player, which basically traverses a "game tree" (tree of possible moves for both players) and quantitatively evaluates the resulting position. This project primarily focuses on how to efficiently traverse such a game tree and create an effective heuristic to maximize the strength of a chess engine. To test the engine's strength, we have scripts that allow the engine to play against humans or Stockfish (a top open source chess engine) at different rating levels. Our engine outperformed our expectations, and has won games against the 2400 rated Stockfish engine, which basically means this engine is strong enough to be a titled player in the chess world.

Background Information

Game Tree Traversal

Minimax

The most basic tree traversal is a minimax algorithm. This algorithm assumes that white and black will play their optimal moves and tries to find the most favorable move *assuming that both players play perfectly*. Essentially it is a depth first search at a fixed depth, and once that depth is reached, a static evaluation of the position is calculated. More information can be found [here](#). While the space complexity of this is negligible, the time complexity is m^p , where each player has roughly m moves per ply p (a ply is half of a move or just one turn for one player). The proof for this is explained in the linked research.

Alpha Beta Pruning

A standard minimax traversal can be very slow, so shortcuts are very useful. Alpha beta pruning stops traversing a portion of the game tree that is already known to be bad. For example, if white makes a move that is pretty good, then it doesn't have to go into detail analyzing lines where it loses a piece because that move is already proven inferior. By skipping these inferior lines, the new average runtime is estimated to be $\sqrt{m^p}$, which can be seen in [this research](#).

ProbCut

This algorithm takes Alpha Beta pruning to the next level by avoiding portions of the gamespace that “probably” won’t be the best line. For example, if a good move has already been found and the current move being observed seems worse *but isn’t proven to be worse than other moves*, ProbCut can still choose to stop checking that move. This makes ProbCut a lot faster, but when it’s done at the same depth as Alpha Beta pruning, it may perform slightly worse. The benefit, though, is that faster algorithms can go much deeper in their searches in the same amount of time, so a faster search (even if it’s less accurate) can actually be a lot better. More info about the ProbCut algorithm can be found [here](#).

Positional and Material Heuristic

Our evaluation function starts with the base evaluation of a board from the perspective of white (which is always the same since board always starts the same) and then for every move being tested (for both white and black) is added to an evaluation stack. The change in the evaluation value for every move is calculated by adding the difference of values for the position of the moved piece, and then adding any extra changes given by pieces being taken.

As mentioned above, we have a set of tables that represent weighting for each piece in each position on the board. This represents how good this position is for a given piece.

Our Implementation

Our project involved creating a python script that plays chess against itself, Stockfish, and human opponents. We created a custom class for storing a board and its heuristic updated in real time, as well as an engine class that can traverse the tree using any implemented algorithm at a custom depth.

What we can do

We built an algorithm that uses minimax and another algorithm that uses alpha beta pruning for move prediction. A minimax approach of depth 3 is as fast as an alpha-beta approach of depth 4, which shows how our alpha-beta approach is significantly faster.

We have a heuristic that constantly updates when a new move is made. Because our heuristic is purely focused on material and positioning of pieces, this allows us to use the existing heuristic and only change the moved piece’s evaluation when updating the new heuristic. This saves the algorithm a lot of time.

Our algorithm has played a variety of games, many of which are stored in the games / folder of the project repository as PGN files. It has beat Stockfish in a variety of games as shown below.

Results

We ran a set of test games against Stockfish with different depths of search in our custom algorithm, we then took these and calculated (a) the average the win percentage of our algorithm and (b) the average time it took for the whole game to play out. You can see these two graphs below.

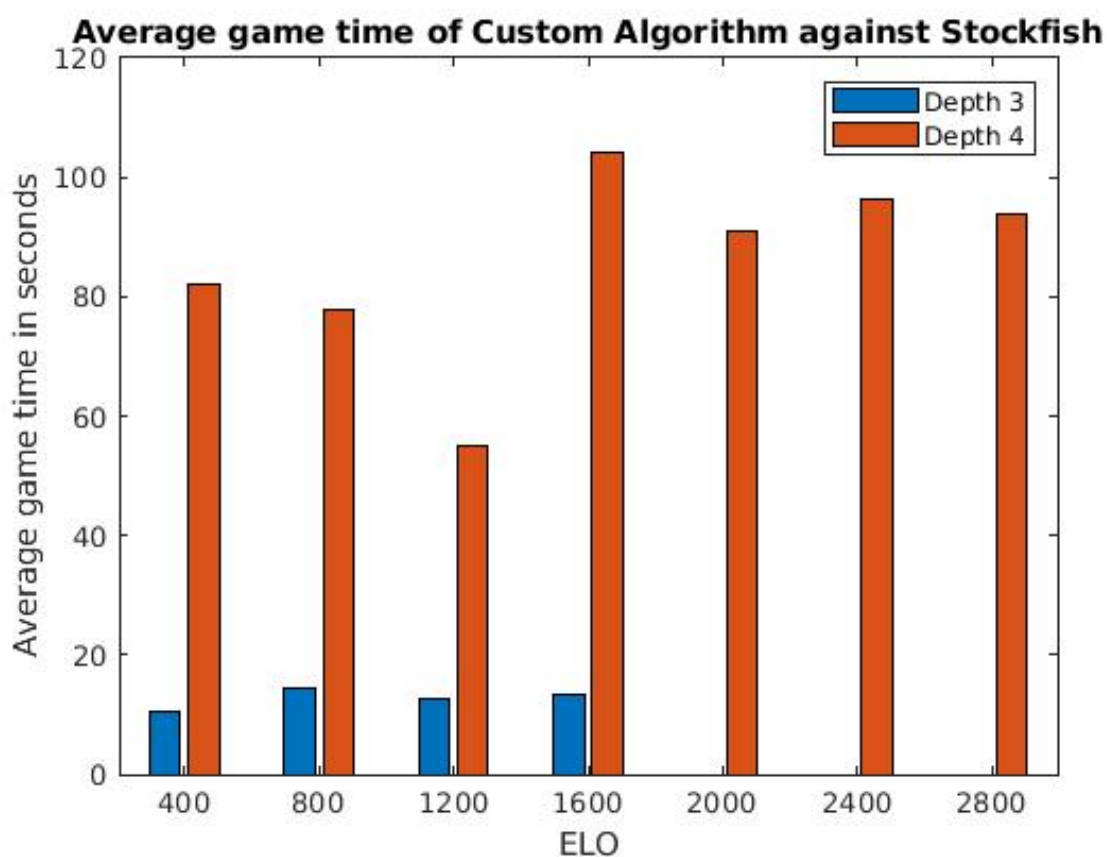


Figure 1: Gametime Graph

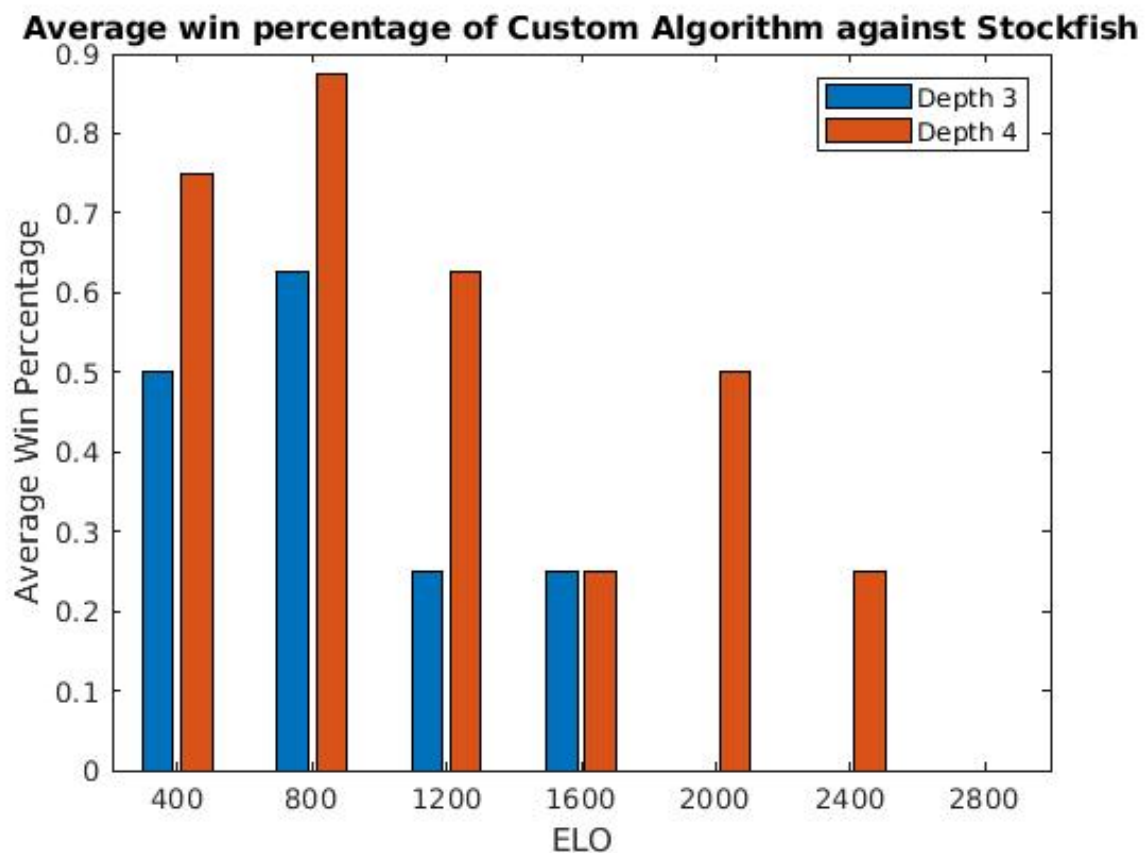


Figure 2: Win Percentage Graph

As shown, the depth 3 engine moved almost immediately. Given that the games had about 70 moves each, each move took a fraction of a second. However, it took about one second per move with the depth 4 engine. The depth 5 engine was so slow (it took anywhere from 30s to a couple of minutes per move) that it did not make sense to even put it on this scale, and we decided that the depth 4 engine was the strongest engine that ran within a reasonable time.

The engine performed significantly better at depth 4, faring well against even 2000-2400 rated opponents, while the depth 3 engine was still decent against a 1600 opponent. Further analysis of these games shows that while the engine is relatively solid, a fixed BFS doesn't know how to win an endgame where there may be a forced checkmate in, say, 10 moves, but it just doesn't have the intuition to see this checkmate so many moves ahead. These can cause some obviously winning positions to end up drawn for the engine. Overall, both of these engines were pretty decent.

As a simple test, Andrew (~1650 Elo) played the depth 4 engine and reached a fairly equal endgame which was basically drawn. The engine seemed to play solid moves and was a decent opponent.

Our presentation includes GIFs of some of the games our engine played, and the [lichess study here](#)

has records of games played with depths 4-5 against Stockfish ELO 1600-3200. All of the .pgn files (chess notation files) from our study are in the games / folder, or you can create a new game by running showdown .py with custom settings for engine depth and other information.

Conclusions

All in all, we are very happy with the engine we were created. It is not perfect by any means, but it is a very decent engine that holds up against other high level engines. One of the big things we could do to improve our engine is make its endgame choices better. We could try and implement an endgame tablebase which introduces a new heuristic and algorithm customized for endgame play. We can add features like ProbCut to make the current search algorithm more efficient. Finally, an opening book would introduce a third heuristic and list of precalculated opening moves that reach more desirable middlegames, since the all-purpose heuristic we implemented is best suited for middlegames and not good but not ideal for opening play.